

```

worksheet7-sambhav.ipynb ⋆ ⓘ
File Edit View Insert Runtime Tools Help
Commands + Code + Text ▶ Run all ▶

Open in Colab

[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

[2]: from sklearn.datasets import fetch_california_housing, load_breast_cancer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression, Ridge, Lasso, LogisticRegression
from sklearn.metrics import mean_squared_error, accuracy_score

[3]: # PART 1: REGRESSION (California Housing)
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split

# Generate synthetic data with similar characteristics
X, y = make_regression(n_samples=20000, n_features=8, noise=10, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

[4]: # Baseline Linear Regression (No Regularization)
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

# Predictions
y_train_pred = lin_reg.predict(X_train)
y_test_pred = lin_reg.predict(X_test)

# MSE
train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_test_pred)

train_mse, test_mse

[5]: (99.78476792958459, 100.88809768997487)

[6]: # Step 2: Hyperparameter Tuning (Ridge & Lasso)
alpha_grid = {'alpha': [0.01, 0.1, 1, 10, 100]}
# Ridge Regression (L2)
ridge_cv = GridSearchCV(
    Ridge(), alpha_grid, cv=5, scoring='neg_mean_squared_error'
)

```

```

worksheet7-sambhav.ipynb ⋆ ⓘ
File Edit View Insert Runtime Tools Help
Commands + Code + Text ▶ Run all ▶

[1]: ridge_cv.fit(X_train, y_train)

ridge_cv.best_params_, ridge_cv.best_mse

[2]: ({'alpha': 0.1}, 100.88851078915167)

[3]: # Lasso Regression (L1)
lasso = Lasso(max_iter=10000)
lasso_cv = GridSearchCV(
    lasso, alpha_grid, cv=5, scoring='neg_mean_squared_error'
)
lasso_cv.fit(X_train, y_train)

best_lasso = lasso_cv.best_estimator_

[4]: best_lasso

[5]: lasso_cv.best_params_, lasso_cv.best_mse

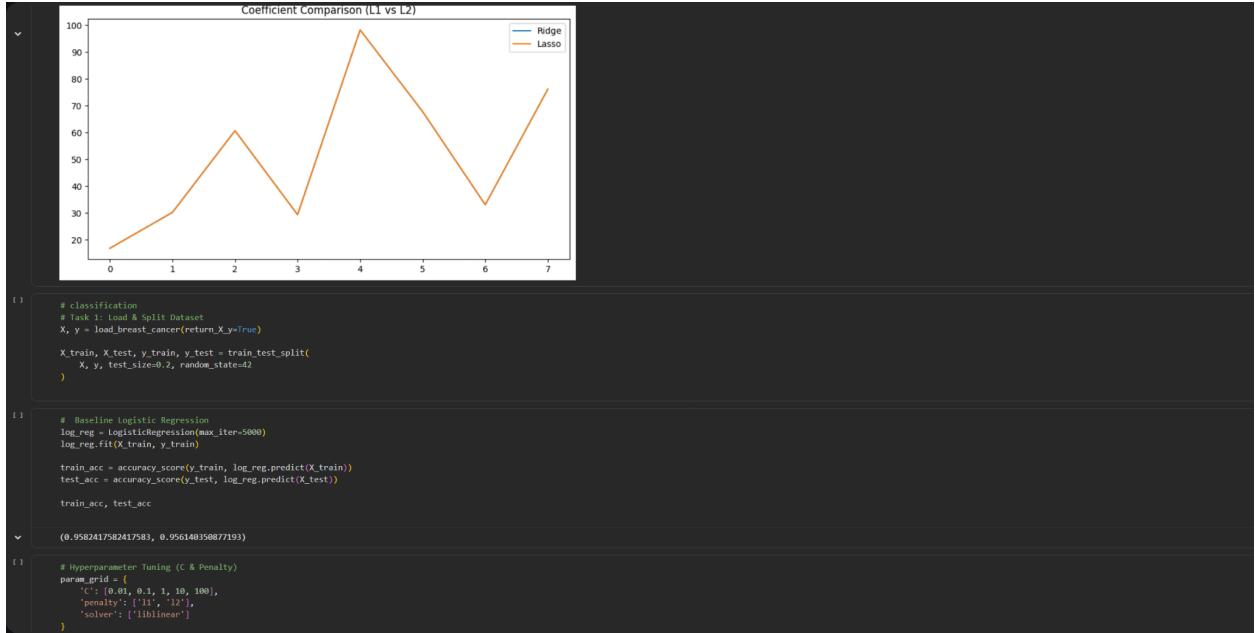
[6]: ({'alpha': 0.01}, 100.994938666345)

[7]: # L1 vs L2 Comparison
pd.DataFrame({
    "Model": ["Linear", "Ridge", "Lasso"],
    "Train MSE": [train_mse, ridge_cv.best_mse, lasso_cv.best_mse],
    "Test MSE": [test_mse, ridge_cv.best_mse, lasso_cv.best_mse]
})

[8]: Model Train MSE Test MSE
0 Linear 99.784768 100.888098
1 Ridge 99.784769 100.888511
2 Lasso 99.785560 100.904904

[9]: # Coefficient Comparison
plt.figure(figsize=(10,5))
plt.plot(best_ridge.coef_, label='Ridge')
plt.plot(best_lasso.coef_, label='Lasso')
plt.legend()
plt.title("Coefficient Comparison (L1 vs L2)")
plt.show()

```



```

() # Hyperparameter Tuning (C & Penalty)
param_grid = (
    'C': [0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear']
)

() log_cv = GridSearchCV(
    LogisticRegression(max_iter=5000),
    param_grid,
    cv=5,
    scoring='accuracy'
)
log_cv.fit(X_train, y_train)
best_log = log_cv.best_estimator_

() best_log, accuracy_score(y_test, best_log.predict(X_test))

()
(LogisticRegression(C=100, max_iter=5000, penalty='l1', solver='liblinear'),
 0.982456140508771)

()
# L1 vs L2 Logistic Regression
log_l1 = LogisticRegression(
    penalty='l1', C=log_cv.best_params_['C'],
    solver='liblinear', max_iter=5000
)
log_l2 = LogisticRegression(
    penalty='l2', C=log_cv.best_params_['C'],
    solver='liblinear', max_iter=5000
)
log_l1.fit(X_train, y_train)
log_l2.fit(X_train, y_train)

()
LogisticRegression
LogisticRegression(C=100, max_iter=5000, solver='liblinear')

()
pd.DataFrame({
    "Model": ["Logistic L1", "Logistic L2"],
    "Train Accuracy": [
        accuracy_score(y_train, log_l1.predict(X_train)),
        accuracy_score(y_train, log_l2.predict(X_train))
    ],
    "Test Accuracy": [
        accuracy_score(y_test, log_l1.predict(X_test)),
        accuracy_score(y_test, log_l2.predict(X_test))
    ]
})

()
pd.DataFrame({
    "Model": ["Logistic L1", "Logistic L2"],
    "Train Accuracy": [
        accuracy_score(y_train, log_l1.predict(X_train)),
        accuracy_score(y_train, log_l2.predict(X_train))
    ],
    "Test Accuracy": [
        accuracy_score(y_test, log_l1.predict(X_test)),
        accuracy_score(y_test, log_l2.predict(X_test))
    ]
})

()
... Model Train Accuracy Test Accuracy
0 Logistic L1 0.989011 0.982456
1 Logistic L2 0.969231 0.956140

()
np.sum(log_l1.coef_ == 0), np.sum(log_l2.coef_ == 0)

()
(np.int64(9), np.int64(0))

()
# Bias-Variance Tradeoff
# Regularization reduces overfitting
# L1: Sparse, interpretable, feature selection
# L2: Smooth, stable, better when many features matter
# Strong regularization → high bias

```